


SEJONG RCV Winter School

- Object Detection(SSD) -


Object Detection

Classification

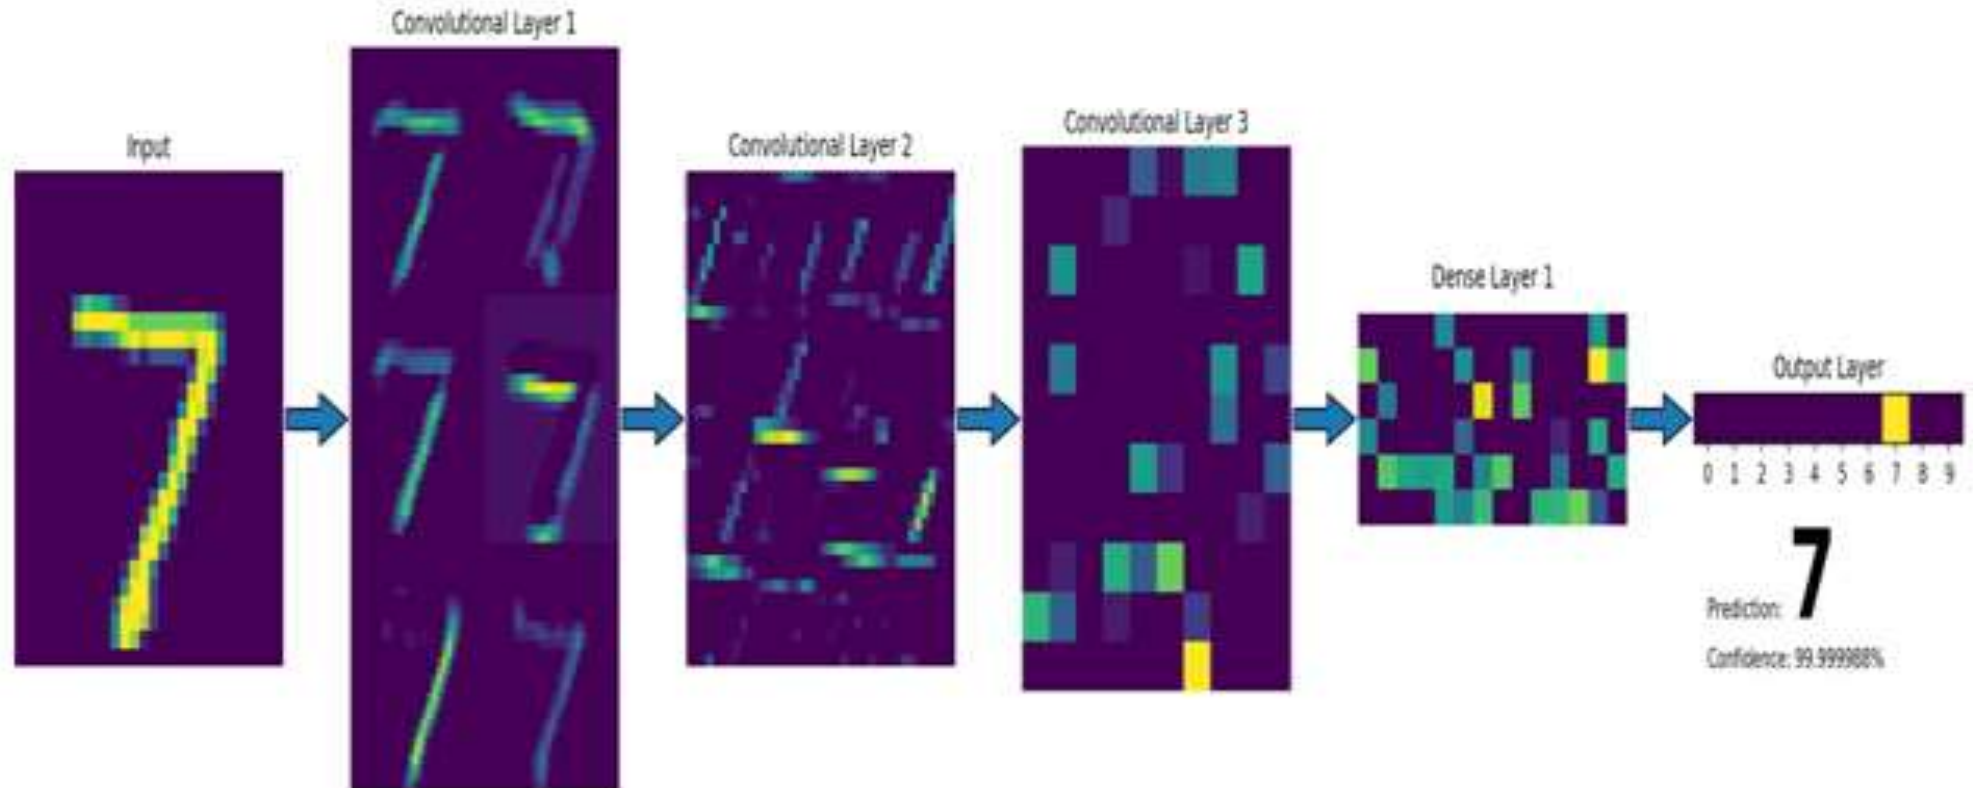
Machine Learning



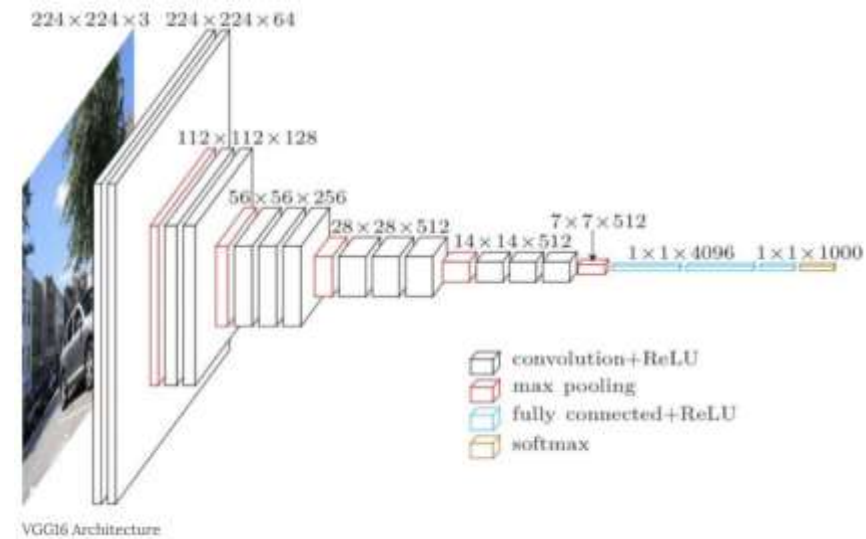
Deep Learning




Classification




Classification



VGG16Net




ResNet



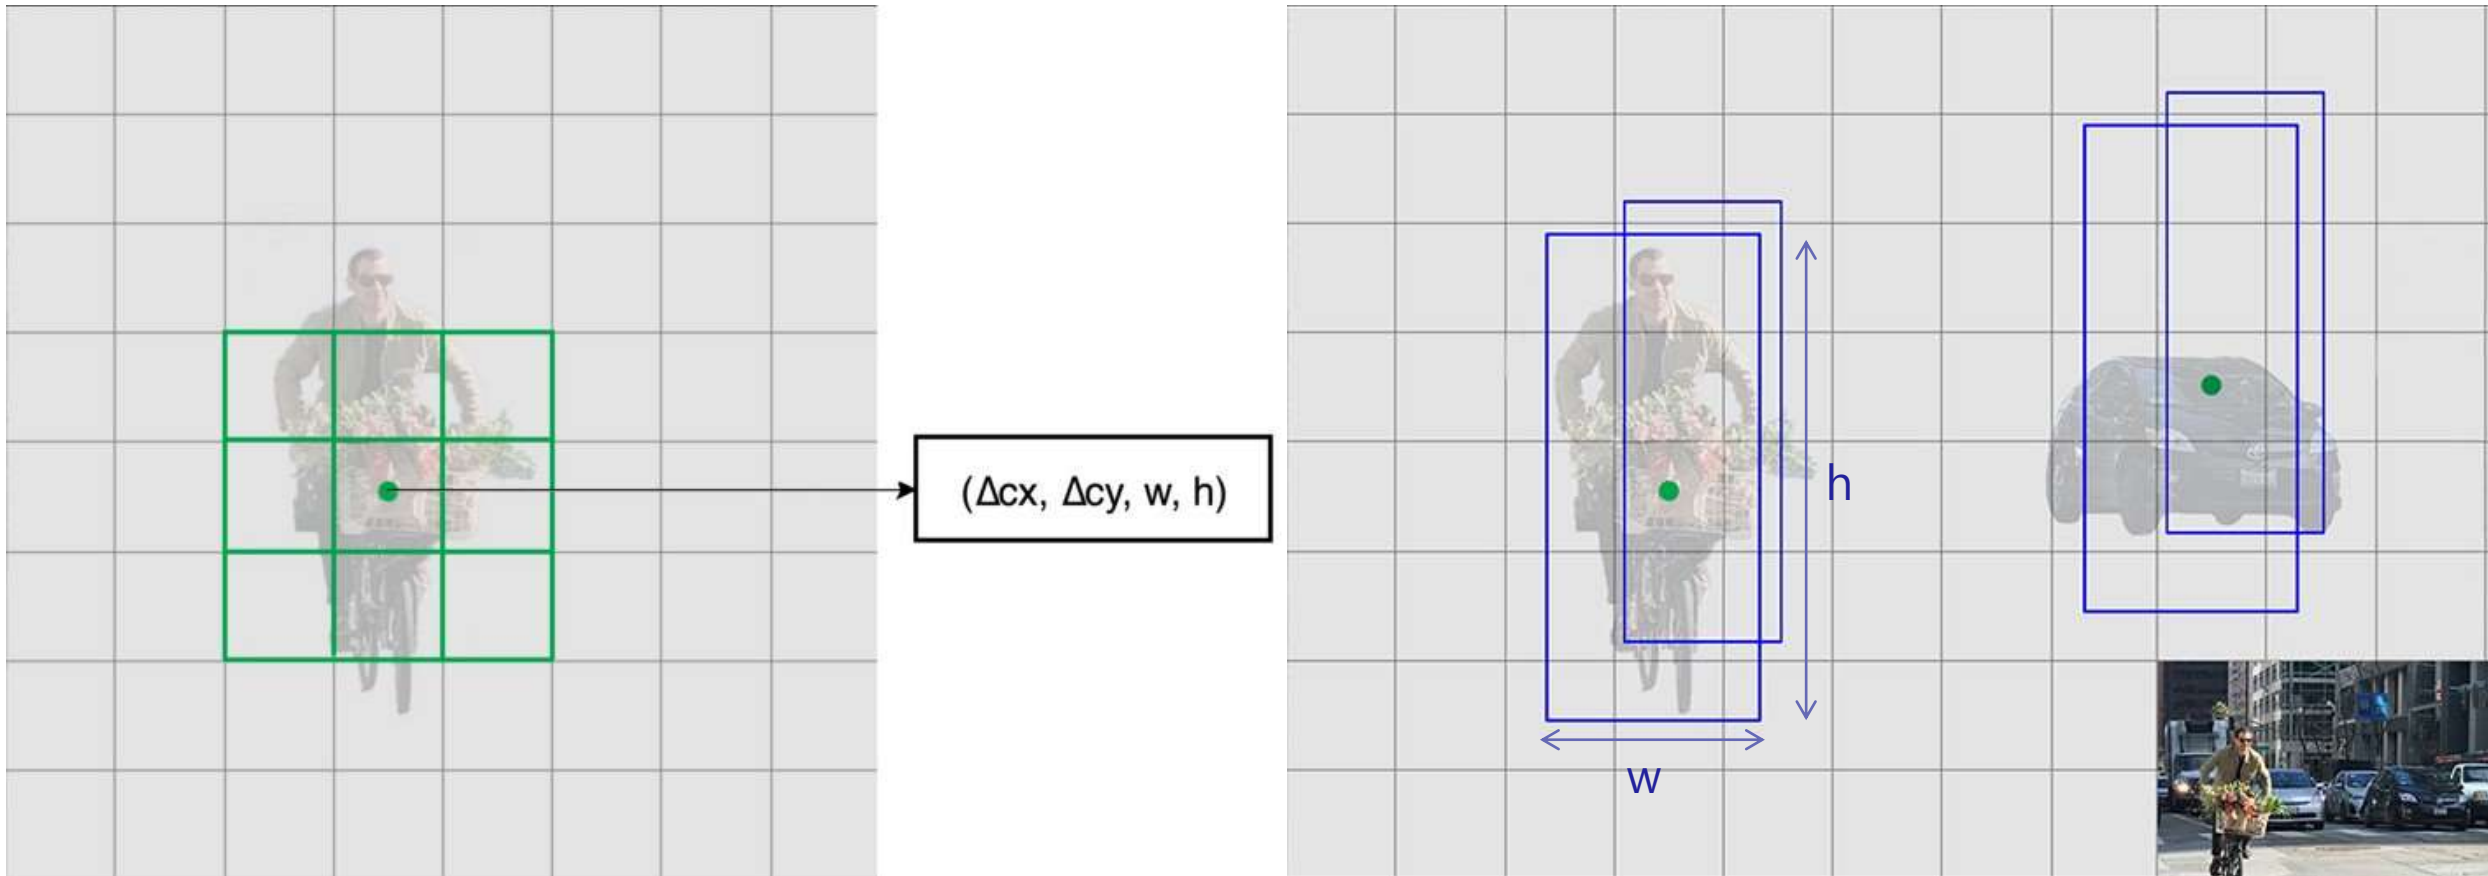
AlexNet

Localization




<https://mc.ai/ssd-object-detection-single-shot-multibox-detector-for-real-time-processing/>

Localization



<https://mc.ai/ssd-object-detection-single-shot-multibox-detector-for-real-time-processing/>

Localization



When priors at a location overshoot the edges of the feature map, they are clipped

R-CNN: Pipeline

R-CNN in detail


- Bounding box regression*
 - Proposal $P = (P_x, P_y, P_w, P_h)$ / Ground-truth $G = (G_x, G_y, G_w, G_h)$
 - For box-scale invariance**, regression target t for the training pair (P, G)

$$t_x = (G_x - P_x) / P_w$$

$$t_y = (G_y - P_y) / P_h$$

$$t_w = \log(G_w / P_w)$$

$$t_h = \log(G_h / P_h)$$



$$\textcircled{1} \quad G = (57.4, 141.0, 97.2, 237.0)$$

$$P = (61.6, 120.0, 107.3, 261.7)$$

$$t_x = (57.4 - 61.6) / 107.3 = -0.039$$

$$t_y = (141.0 - 120.0) / 261.7 = 0.080$$

$$t_w = \log(97.2 / 107.3) = -0.043$$

$$t_h = \log(237.0 / 261.7) = -0.043$$

$$\textcircled{2} \quad G = (408.0, 167.0, 29.9, 73.0)$$

$$P = (405.8, 170.8, 31.5, 76.9)$$

$$t_x = (408.0 - 405.8) / 31.5 = 0.070$$

$$t_y = (167.0 - 170.8) / 76.9 = -0.049$$

$$t_w = \log(29.9 / 31.5) = -0.023$$

$$t_h = \log(73.0 / 76.9) = -0.023$$

SSD(Single Shot MultiBox Detector)

<https://mc.ai/ssd-object-detection-single-shot-multibox-detector-for-real-time-processing/>

SSD: Single Shot MultiBox Detector

Wei Liu¹, Dragomir Anguelov², Dumitru Erhan³, Christian Szegedy³,
Scott Reed⁴, Cheng-Yang Fu¹, Alexander C. Berg¹

¹UNC Chapel Hill ²Zoox Inc. ³Google Inc. ⁴University of Michigan, Ann-Arbor

¹wliu@cs.unc.edu, ²drago@zoox.com, ³{dumitru,szegedy}@google.com,

⁴reedscot@umich.edu, ¹{cyfu,aberg}@cs.unc.edu

Abstract. We present a method for detecting objects in images using a single deep neural network. Our approach, named SSD, discretizes the output space of bounding boxes into a set of default boxes over different aspect ratios and scales per feature map location. At prediction time, the network generates scores for the presence of each object category in each default box and produces adjustments to the box to better match the object shape. Additionally, the network combines predictions from multiple feature maps with different resolutions to naturally handle objects of various sizes. SSD is simple relative to methods that require object proposals because it completely eliminates proposal generation and subsequent pixel or feature resampling stages and encapsulates all computation in a single network. This makes SSD easy to train and straightforward to integrate into systems that require a detection component. Experimental results on the PASCAL VOC, COCO, and ILSVRC datasets confirm that SSD has competitive accuracy to methods that utilize an additional object proposal step and is much faster, while providing a unified framework for both training and inference. For 300×300 input, SSD achieves 74.3% mAP¹ on VOC2007 test at 59 FPS on a Nvidia Titan X and for 512×512 input, SSD achieves 76.9% mAP, outperforming a comparable state-of-the-art Faster R-CNN model. Compared to other single stage methods, SSD has much better accuracy even with a smaller input image size. Code is available at: <https://github.com/weiliu89/caffe/tree/ssd>.

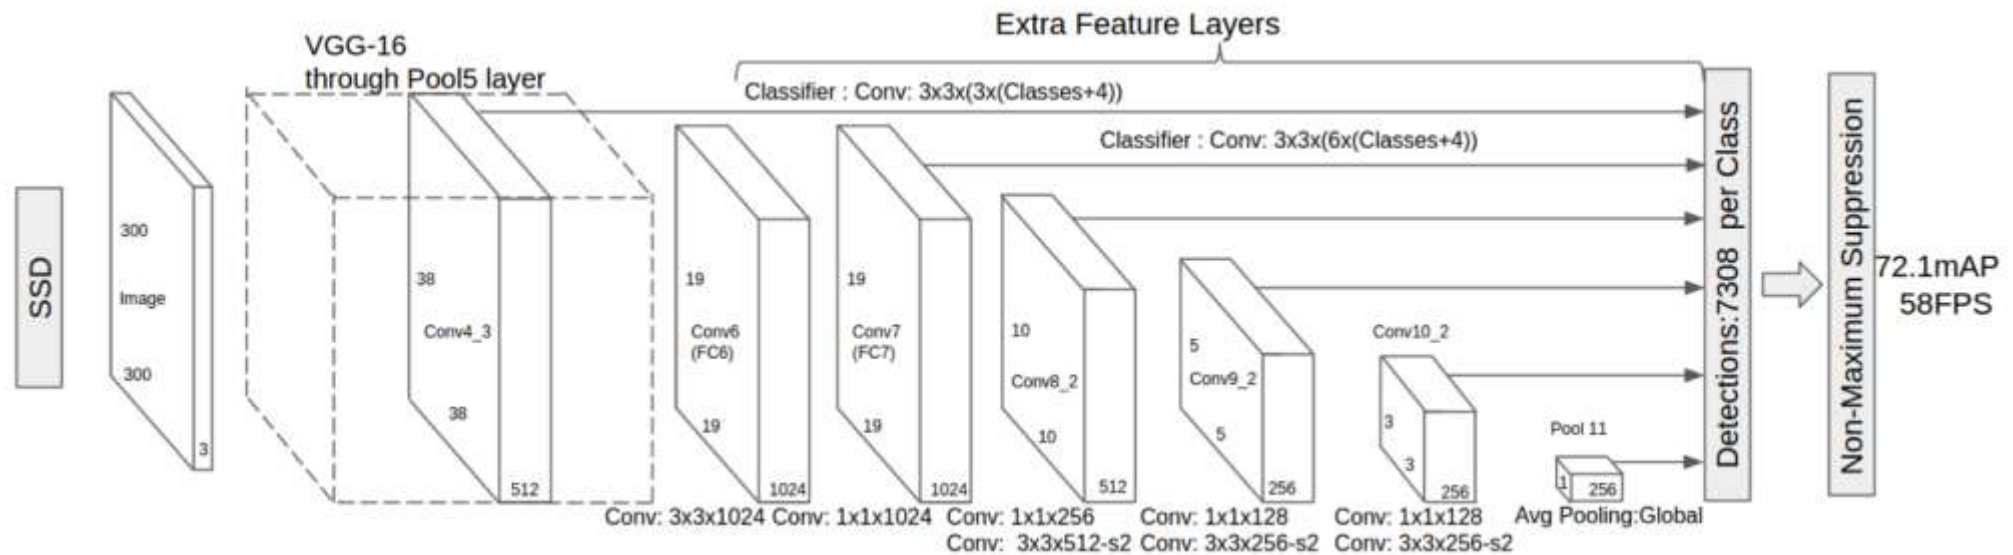
Keywords: Real-time Object Detection; Convolutional Neural Network



Classification + Localization


<https://arxiv.org/pdf/1512.02325.pdf>

Classification + Localization




SSD(Single Shot multi-box Detector)

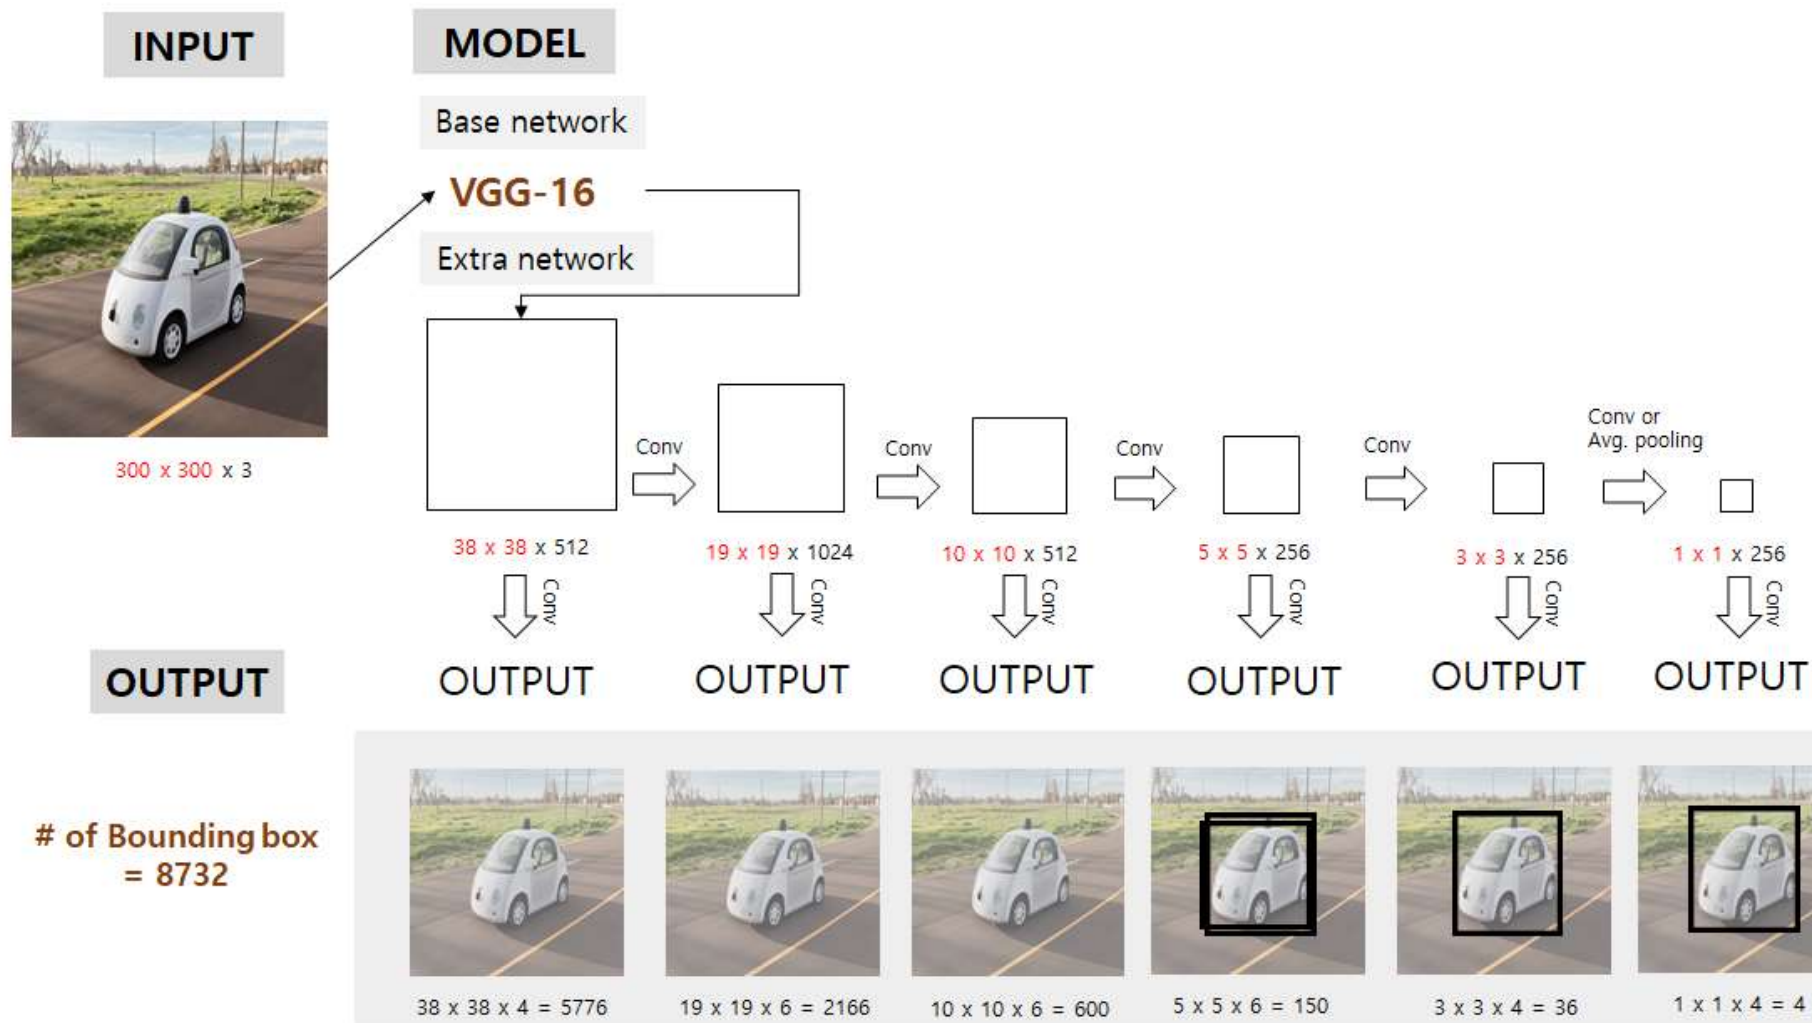
Classification + Localization



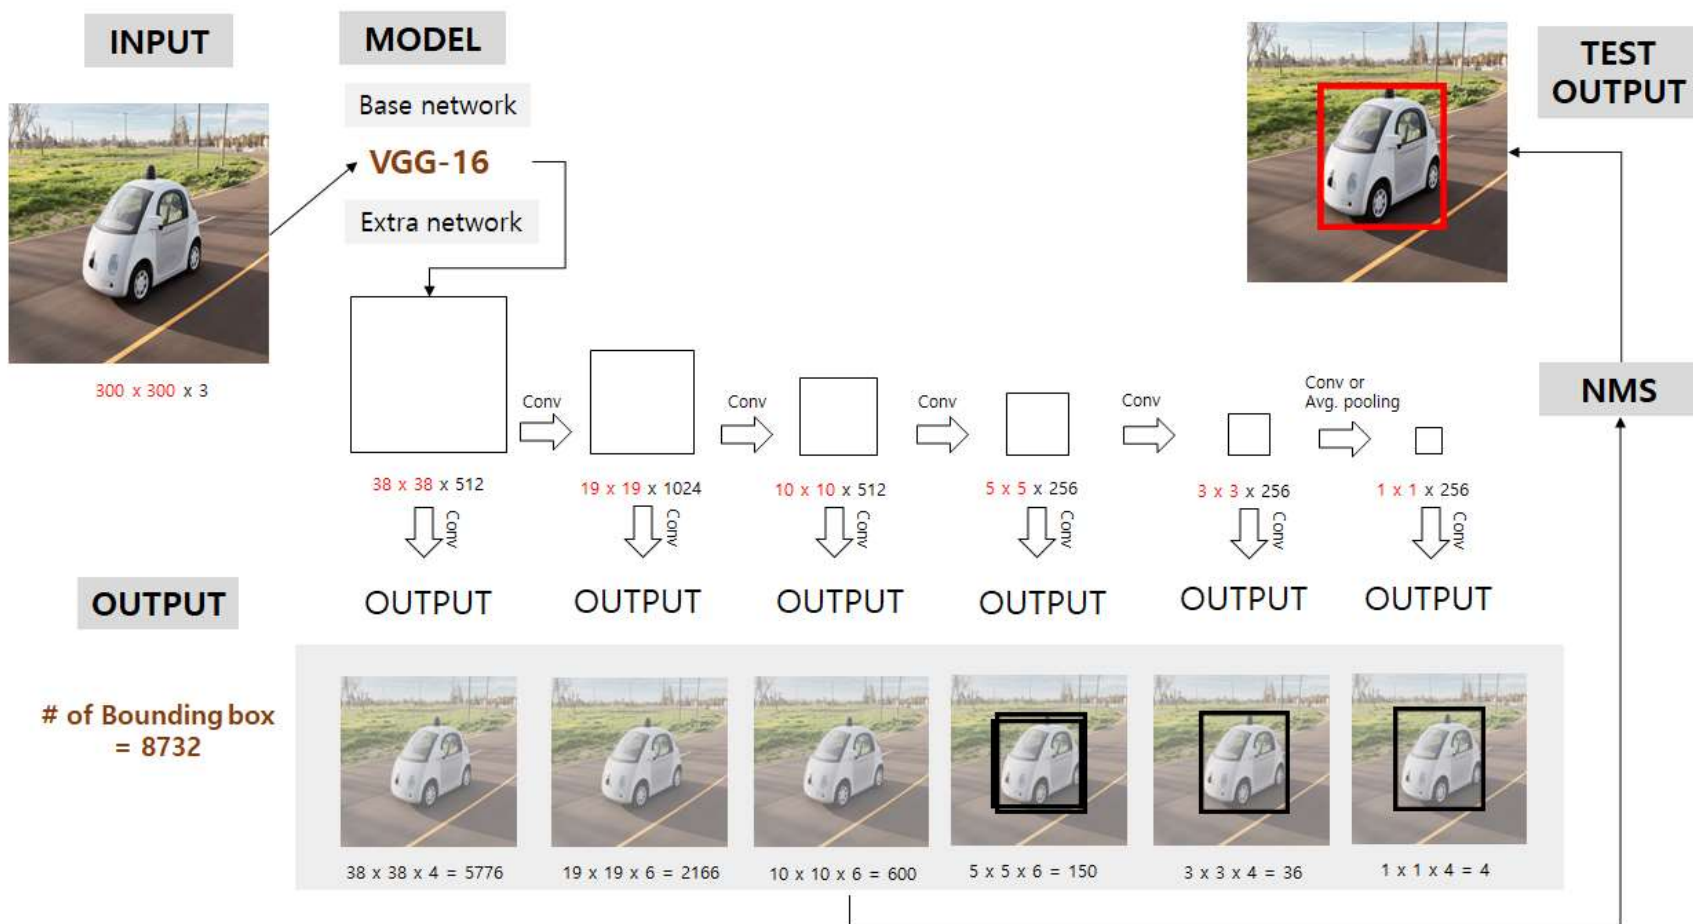
Classification + Localization



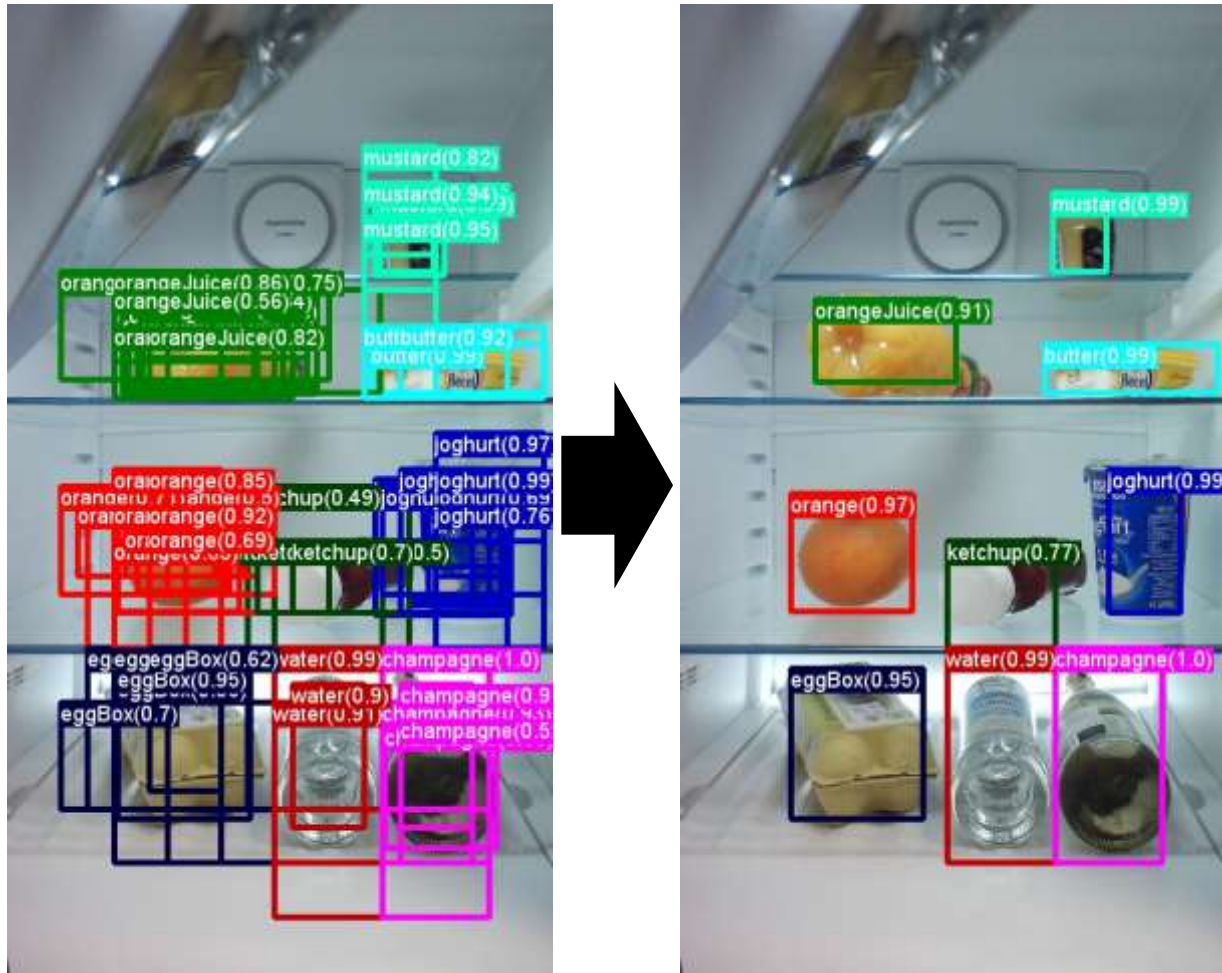
Classification + Localization



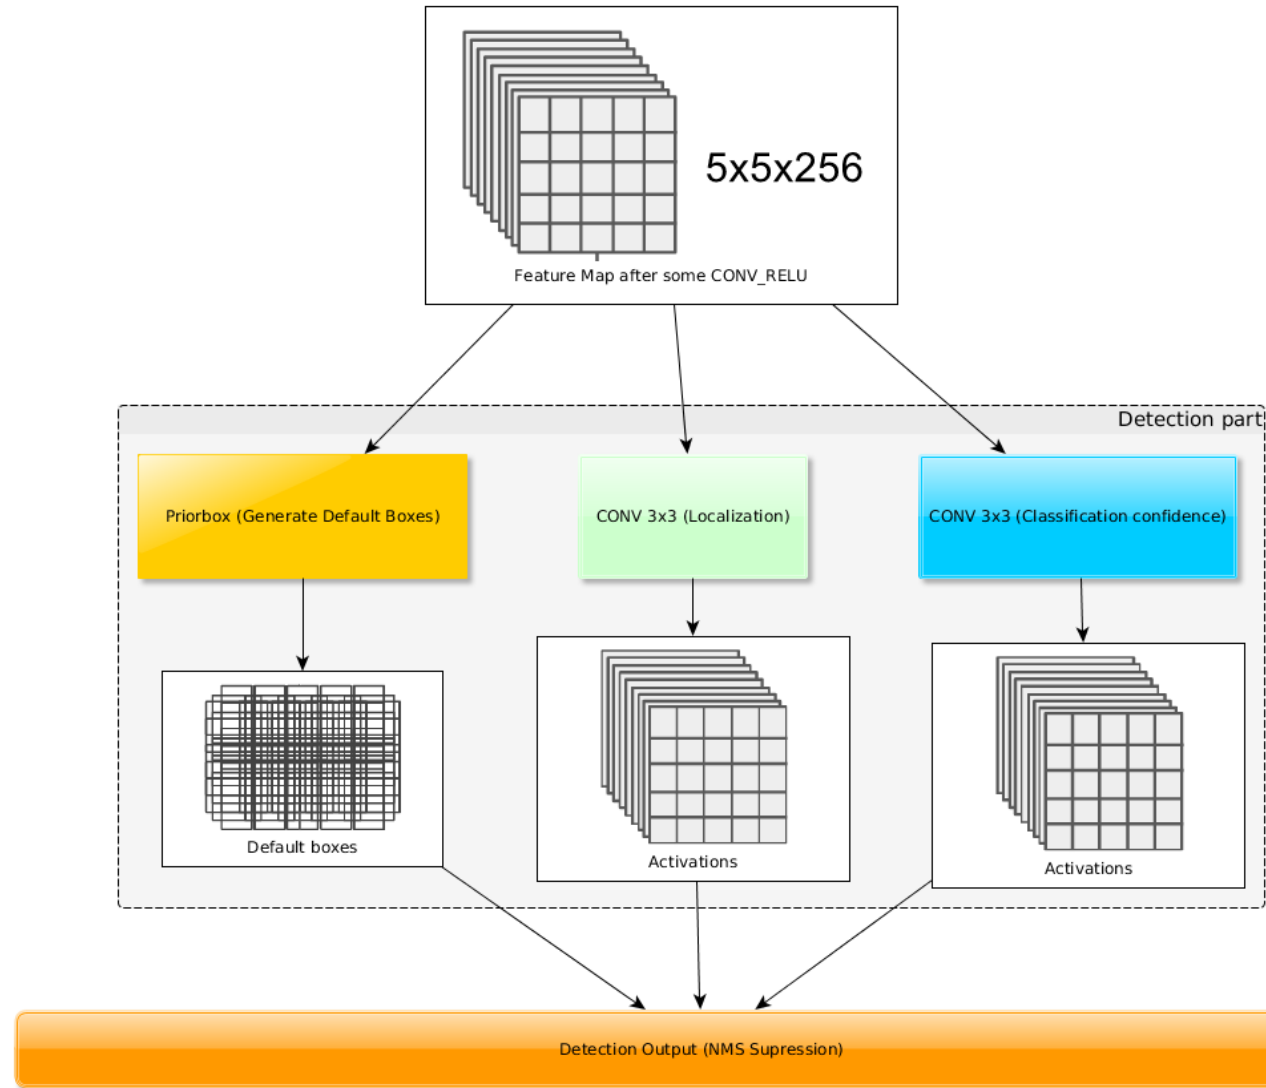
Classification + Localization



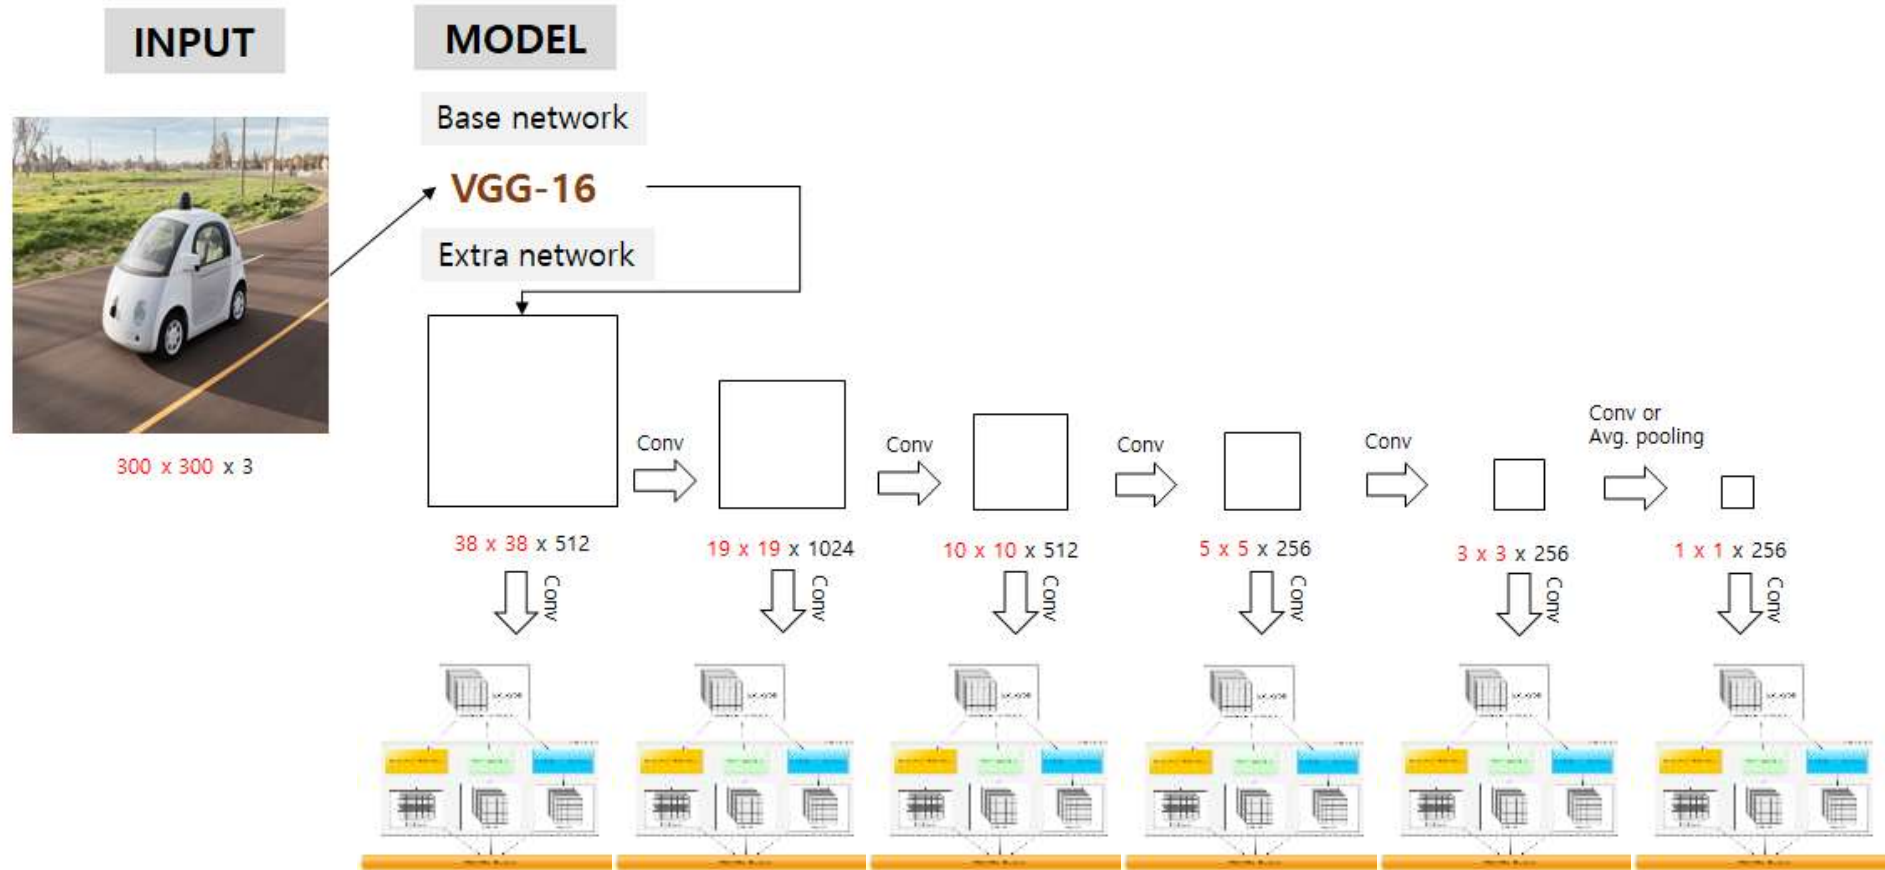
NMS



Classification + Localization



Classification + Localization



sgrvinod / [a-PyTorch-Tutorial-to-Object-Detection](#) Unwatch 36 Unstar 1k Fork 248


[Code](#) [Issues 22](#) [Pull requests 3](#) [Actions](#) [Projects 0](#) [Wiki](#) [Security](#) [Insights](#)










SSD: Single Shot MultiBox Detector | a PyTorch Tutorial to Object Detection

[pytorch](#) [pytorch-tutorial](#) [object-detection](#) [single-shot-multibox-detector](#) [single-shot-detection](#) [object-recognition](#) [ssd](#) [tutorial](#) [detection](#)

51 commits 1 branch 0 packages 0 releases 2 contributors

Branch: master [New pull request](#) [Create new file](#) [Upload files](#) [Find file](#) [Clone or download](#)

 sgrvinod updated model checkpoint (trained without early stopping) Latest commit 9cad25b 14 days ago

 img	updated some images	12 months ago
 README.md	updated model checkpoint (trained without early stopping)	14 days ago
 create_data_lists.py	Initial commit	2 years ago
 datasets.py	added tutorial content	12 months ago
 detect.py	changes	14 days ago
 eval.py	changes	14 days ago
 model.py	added tutorial content	12 months ago
 train.py	changes	14 days ago
 utils.py	changes	14 days ago

<https://github.com/sgrvinod/a-PyTorch-Tutorial-to-Object-Detection>

SEJONG RCV Winter School

- Object Detection(Halfway) -

Multispectral Deep Neural Networks for Pedestrian Detection

Jingjing Liu¹

<http://paul.rutgers.edu/~jl1322>

Shaoting Zhang²

szhang16@uncc.edu

Shu Wang¹

sw498@cs.rutgers.edu

Dimitris N. Metaxas¹


<https://www.cs.rutgers.edu/~dnm>


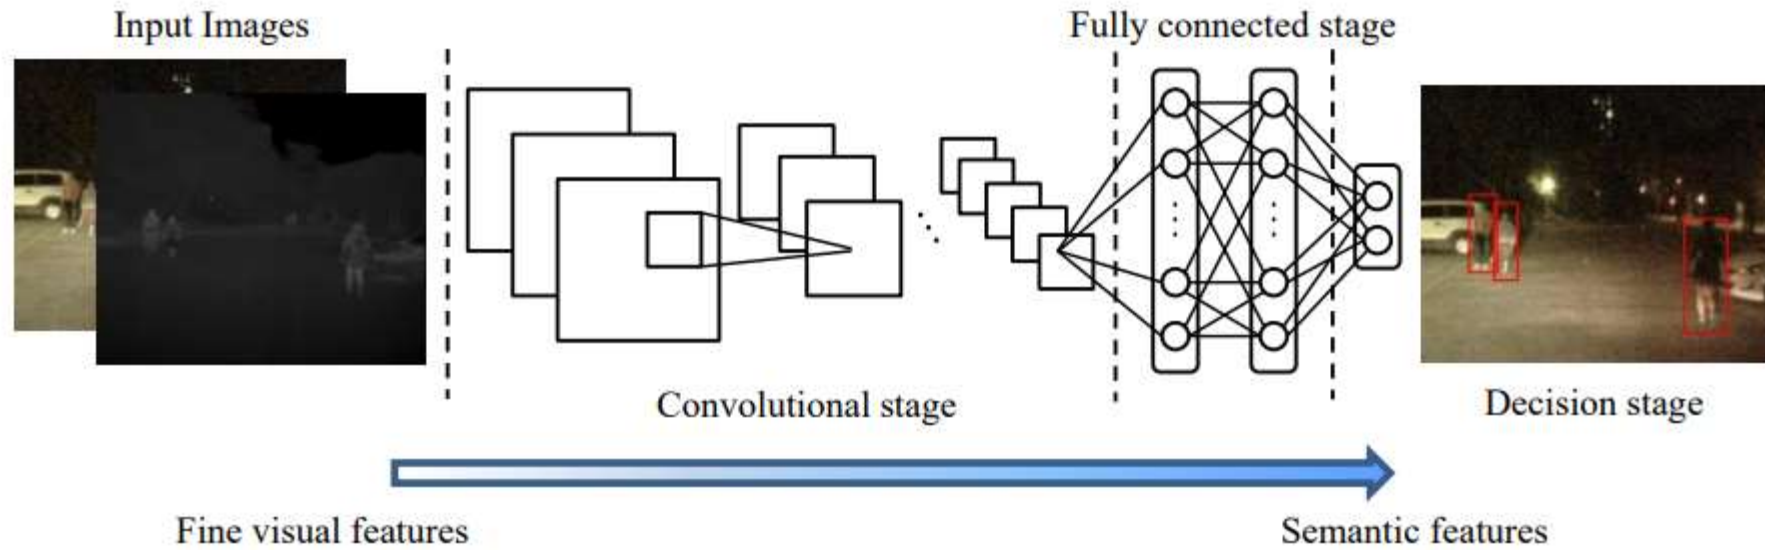
¹ Department of Computer Science
Rutgers University
Piscataway, NJ, USA

² Department of Computer Science
UNC Charlotte
Charlotte, NC, USA


Abstract

Multispectral pedestrian detection is essential for around-the-clock applications, *e.g.*, surveillance and autonomous driving. We deeply analyze Faster R-CNN for multispectral pedestrian detection task and then model it into a convolutional network (ConvNet) fusion problem. Further, we discover that ConvNet-based pedestrian detectors trained by color or thermal images separately provide complementary information in discriminating human instances. Thus there is a large potential to improve pedestrian detection by using color and thermal images in DNNs simultaneously. We carefully design four ConvNet fusion architectures that integrate two-branch ConvNets on different DNNs stages, all of which yield better performance compared with the baseline detector. Our experimental results on KAIST pedestrian benchmark show that the Halfway Fusion model that performs fusion on the middle-level convolutional features outperforms the baseline method by 11% and yields a missing rate 3.5% lower than the other proposed architectures.







Early Fusion



Halfway Fusion



Late Fusion



Score Fusion

SEJONG RCV Winter School

- Object Detection(Static Fusion & Adaptive Fusion) -

Single-Shot Adaptive Fusion Network for Robust Multispectral Pedestrian Detection


*김지원¹, *조 원¹, 남현호¹, 황순민¹, 노치원², 김남일³, 최유경¹


세종대학교 Robotics Computer Vision (RCV) 연구실¹, POTENIT², NAVER LABS³

{jwkim, jwon, hhnam, smhwang, ykchoi}@rcv.sejong.ac.kr, cwroh@potenit.com, namil.kim@naverlabs.com

요약

컴퓨터 비전 분야에서 밤과 낮, 날씨의 변화 등 다양한 자연환경에 대한 환경변화 강인성은 극복 해야할 중요 주제로, 오래전부터 다양한 분야에서 많은 연구들이 진행되었다. 하지만 일정 수준 예측이 가능한 환경변화 강인성 이외에도, 실제 시스템 운영 시에 센서의 파손 및 고장, 먼지/곤충에 의한 센서의 가려짐 등 예상치 못한 상황들이 일어날 수 있고, 알고리즘 학습 시 고려하지 못했던 이러한 변수들은 실제 테스트 중 큰 성능저하를 야기한다. 본 논문에서는 학습과정에서 예상할 수 없었던 변화에 대해서도 강인성을 확보할 수 있는 새로운 학습방법과 입력 영상에 따라 유동적으로 학습 파라미터를 생성하는 모델을 제안한다. 제안하는 방법론을 KAIST multispectral benchmark 데이터 셋 기반으로 다양한 예측 불가능한 상황들을 모사하여 검증하였으며, 예측 불가능한 상황에 강인하게 동작하는 것뿐 아니라, 정상 조건에서도 SOTA의 성능을 보이는 것을 입증하였다.





```
##### RGB #####

out_vis = F.relu(self.conv2_1_bn_vis(self.conv2_1_vis(out_vis)))
out_vis = F.relu(self.conv2_2_bn_vis(self.conv2_2_vis(out_vis)))
out_vis = self.pool2_vis(out_vis)

out_vis = F.relu(self.conv3_1_bn_vis(self.conv3_1_vis(out_vis)))
out_vis = F.relu(self.conv3_2_bn_vis(self.conv3_2_vis(out_vis)))
out_vis = F.relu(self.conv3_3_bn_vis(self.conv3_3_vis(out_vis)))

out_vis = F.relu(self.conv1x1_sf_vis(out_vis))
#out_vis = self.fusion_R(conv1_2_feats_vis, conv1_2_feats_lwir, out_vis)

##### Thermal #####


out_lwir = F.relu(self.conv2_1_bn_lwir(self.conv2_1_lwir(out_lwir)))
out_lwir = F.relu(self.conv2_2_bn_lwir(self.conv2_2_lwir(out_lwir)))
out_lwir = self.pool2_lwir(out_lwir)

out_lwir = F.relu(self.conv3_1_bn_lwir(self.conv3_1_lwir(out_lwir)))
out_lwir = F.relu(self.conv3_2_bn_lwir(self.conv3_2_lwir(out_lwir)))
out_lwir = F.relu(self.conv3_3_bn_lwir(self.conv3_3_lwir(out_lwir)))

out_lwir = F.relu(self.conv1x1_sf_lwir(out_lwir))
#out_lwir = self.fusion_T(conv1_2_feats_vis, conv1_2_feats_lwir, out_lwir)

out = out_vis*0.5 + out_lwir*0.5

#####
```



```
##### RGB #####
```

```
out_vis = F.relu(self.conv2_1_bn_vis(self.conv2_1_vis(out_vis)))
out_vis = F.relu(self.conv2_2_bn_vis(self.conv2_2_vis(out_vis)))
out_vis = self.pool2_vis(out_vis)
```

```
out_vis = F.relu(self.conv3_1_bn_vis(self.conv3_1_vis(out_vis)))
out_vis = F.relu(self.conv3_2_bn_vis(self.conv3_2_vis(out_vis)))
out_vis = F.relu(self.conv3_3_bn_vis(self.conv3_3_vis(out_vis)))
```

```
# out_vis = F.relu(self.conv1x1_sf_vis(out_vis))
out_vis = self.fusion_R(conv1_2_feats_vis, conv1_2_feats_lwir, out_vis)
```

```
#####
```

```
##### Thermal #####
```

```
out_lwir = F.relu(self.conv2_1_bn_lwir(self.conv2_1_lwir(out_lwir)))
out_lwir = F.relu(self.conv2_2_bn_lwir(self.conv2_2_lwir(out_lwir)))
out_lwir = self.pool2_lwir(out_lwir)
```

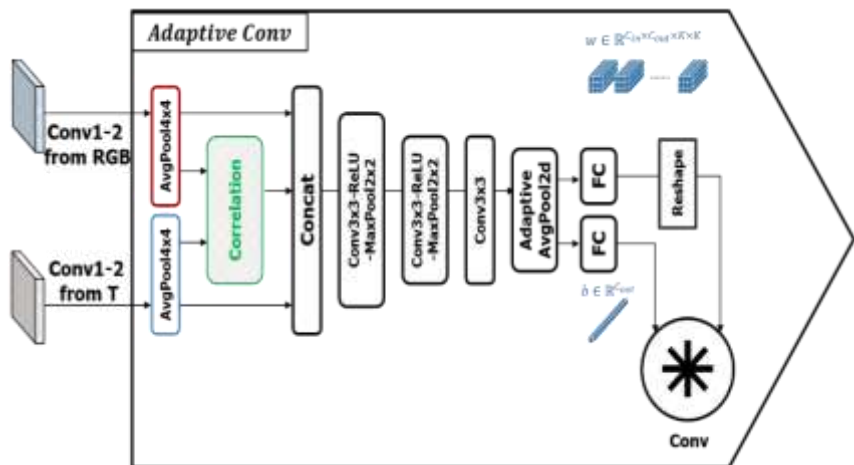
```
out_lwir = F.relu(self.conv3_1_bn_lwir(self.conv3_1_lwir(out_lwir)))
out_lwir = F.relu(self.conv3_2_bn_lwir(self.conv3_2_lwir(out_lwir)))
out_lwir = F.relu(self.conv3_3_bn_lwir(self.conv3_3_lwir(out_lwir)))
```

```
# out_lwir = F.relu(self.conv1x1_sf_lwir(out_lwir))
out_lwir = self.fusion_T(conv1_2_feats_vis, conv1_2_feats_lwir, out_lwir)
```

```
#####
```

```
out = out_vis*0.5 + out_lwir*0.5
```

```
#####
```



```
def forward(self, x, y, hx):

    bs = x.size(0)

    _x = F.avg_pool2d(x.detach(), kernel_size=4, stride=4)
    _y = F.avg_pool2d(y.detach(), kernel_size=4, stride=4)

    c = self.corr( _x, _y )
    c = c.view(c.size(0),c.size(1)*c.size(2),c.size(3),c.size(4))

    feat = t.cat( [_x, _y, c], 1 )

    feat = F.adaptive_avg_pool2d( self.conv(feat), (1, 1) )
    feat = self.conv_post( feat )

    feat = feat.view( bs, -1 )

    xx = list()

    wx = self.wx( feat ).view( bs, self.hidden_ch, self.out_ch, self.kernel_size, self.kernel_size)
    bx = self.bx( feat ).view( bs, self.out_ch)

    for ii in range(bs):
        xx.append( F.conv2d(hx[ii:ii+1,...], wx[ii,...], bx[ii,...], padding=self.padding) )

    x = t.cat( xx, 0 )

    return x
```

```
class AdaptiveFusion(nn.Module):
    def __init__(self, in_ch, hidden_ch, out_ch, kernel_size,patch_size):

        super(AdaptiveFusion, self).__init__()

        self.patch_size = patch_size

        self.in_ch = in_ch
        self.hidden_ch = hidden_ch
        self.out_ch = out_ch

        self.kernel_size = kernel_size
        self.padding = math.floor(kernel_size/2)

        self.corr = SpatialCorrelationSampler(kernel_size=1,patch_size=self.patch_size,stride=1,padding=0,dilation_patch=2)

        self.conv = nn.Sequential(
            nn.Conv2d( in_ch*2 + (self.patch_size)*(self.patch_size) , hidden_ch, kernel_size=3, padding=1, stride=1, bias=False),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d( hidden_ch, hidden_ch, kernel_size=3, padding=1, stride=1, bias=False),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            nn.Conv2d( hidden_ch, hidden_ch, kernel_size=3, padding=1, stride=1, bias=False),
        )

        self.conv_post = nn.Conv2d( hidden_ch, hidden_ch, kernel_size=1, padding=0, stride=1)

        self.wx = nn.Linear( hidden_ch, hidden_ch*out_ch*kernel_size*kernel_size )
        self.bx = nn.Linear( hidden_ch, out_ch )

        self.reset_parameters()
```